

Boston House Data Prediction using popular Regression Algorithms



Importing necessary Libraries

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_boston
import seaborn as sns
%matplotlib inline
```

Loading the Dataset

In [2]:

```
data = load_boston()
data.keys()
```

C:\Users\Shobhandeb\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function load_boston is deprecated; `load_boston` is deprecated in 1.0 and will be removed in 1.2.

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np
```

```
data_url = "http://lib.stat.cmu.edu/datasets/boston"
```

```
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch_california_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

```
for the Ames housing dataset.
warnings.warn(msg, category=FutureWarning)
```

Out[2]:

```
dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename', 'data_module'])
```

In [3]:

```
print(data.DESCR)
```

```
.. _boston_dataset:
```

```
Boston house prices dataset
-----
```

```
**Data Set Characteristics:**
```

```
:Number of Instances: 506
```

```
:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14)
is usually the target.
```

```
:Attribute Information (in order):
```

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of black people by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

```
:Missing Attribute Values: None
```

```
:Creator: Harrison, D. and Rubinfeld, D.L.
```

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address re

The Boston house price data has been used in many machine learning papers and address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

In [4]:

```
print(data.data)
```

```
[[6.3200e-03 1.8000e+01 2.3100e+00 ... 1.5300e+01 3.9690e+02 4.9800e+00]
 [2.7310e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9690e+02 9.1400e+00]
 [2.7290e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9283e+02 4.0300e+00]
 ...
 [6.0760e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 5.6400e+00]
 [1.0959e-01 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9345e+02 6.4800e+00]
 [4.7410e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 7.8800e+00]]
```

In [5]:

```
print(data.target)
```

```
[24. 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 15. 18.9 21.7 20.4
 18.2 19.9 23.1 17.5 20.2 18.2 13.6 19.6 15.2 14.5 15.6 13.9 16.6 14.8
 18.4 21. 12.7 14.5 13.2 13.1 13.5 18.9 20. 21. 24.7 30.8 34.9 26.6
 25.3 24.7 21.2 19.3 20. 16.6 14.4 19.4 19.7 20.5 25. 23.4 18.9 35.4
 24.7 31.6 23.3 19.6 18.7 16. 22.2 25. 33. 23.5 19.4 22. 17.4 20.9
 24.2 21.7 22.8 23.4 24.1 21.4 20. 20.8 21.2 20.3 28. 23.9 24.8 22.9
 23.9 26.6 22.5 22.2 23.6 28.7 22.6 22. 22.9 25. 20.6 28.4 21.4 38.7
 43.8 33.2 27.5 26.5 18.6 19.3 20.1 19.5 19.5 20.4 19.8 19.4 21.7 22.8
 18.8 18.7 18.5 18.3 21.2 19.2 20.4 19.3 22. 20.3 20.5 17.3 18.8 21.4
 15.7 16.2 18. 14.3 19.2 19.6 23. 18.4 15.6 18.1 17.4 17.1 13.3 17.8
 14. 14.4 13.4 15.6 11.8 13.8 15.6 14.6 17.8 15.4 21.5 19.6 15.3 19.4
 17. 15.6 13.1 41.3 24.3 23.3 27. 50. 50. 50. 22.7 25. 50. 23.8
 23.8 22.3 17.4 19.1 23.1 23.6 22.6 29.4 23.2 24.6 29.9 37.2 39.8 36.2
 37.9 32.5 26.4 29.6 50. 32. 29.8 34.9 37. 30.5 36.4 31.1 29.1 50.
 33.3 30.3 34.6 34.9 32.9 24.1 42.3 48.5 50. 22.6 24.4 22.5 24.4 20.
 21.7 19.3 22.4 28.1 23.7 25. 23.3 28.7 21.5 23. 26.7 21.7 27.5 30.1
 44.8 50. 37.6 31.6 46.7 31.5 24.3 31.7 41.7 48.3 29. 24. 25.1 31.5
 23.7 23.3 22. 20.1 22.2 23.7 17.6 18.5 24.3 20.5 24.5 26.2 24.4 24.8
 29.6 42.8 21.9 20.9 44. 50. 36. 30.1 33.8 43.1 48.8 31. 36.5 22.8
 30.7 50. 43.5 20.7 21.1 25.2 24.4 35.2 32.4 32. 33.2 33.1 29.1 35.1
 45.4 35.4 46. 50. 32.2 22. 20.1 23.2 22.3 24.8 28.5 37.3 27.9 23.9
 21.7 28.6 27.1 20.3 22.5 29. 24.8 22. 26.4 33.1 36.1 28.4 33.4 28.2
 22.8 20.3 16.1 22.1 19.4 21.6 23.8 16.2 17.8 19.8 23.1 21. 23.8 23.1
 20.4 18.5 25. 24.6 23. 22.2 19.3 22.6 19.8 17.1 19.4 22.2 20.7 21.1
 19.5 18.5 20.6 19. 18.7 32.7 16.5 23.9 31.2 17.5 17.2 23.1 24.5 26.6
 22.9 24.1 18.6 30.1 18.2 20.6 17.8 21.7 22.7 22.6 25. 19.9 20.8 16.8
 21.9 27.5 21.9 23.1 50. 50. 50. 50. 50. 13.8 13.8 15. 13.9 13.3
 13.1 10.2 10.4 10.9 11.3 12.3 8.8 7.2 10.5 7.4 10.2 11.5 15.1 23.2
 9.7 13.8 12.7 13.1 12.5 8.5 5. 6.3 5.6 7.2 12.1 8.3 8.5 5.
 11.9 27.9 17.2 27.5 15. 17.2 17.9 16.3 7. 7.2 7.5 10.4 8.8 8.4
 16.7 14.2 20.8 13.4 11.7 8.3 10.2 10.9 11. 9.5 14.5 14.1 16.1 14.3
 11.7 13.4 9.6 8.7 8.4 12.8 10.5 17.1 18.4 15.4 10.8 11.8 14.9 12.6
 14.1 13. 13.4 15.2 16.1 17.8 14.9 14.1 12.7 13.5 14.9 20. 16.4 17.7
 19.5 20.2 21.4 19.9 19. 19.1 19.1 20.1 19.9 19.6 23.2 29.8 13.8 13.3
 16.7 12. 14.6 21.4 23. 23.7 25. 21.8 20.6 21.2 19.1 20.6 15.2 7.
 8.1 13.6 20.1 21.8 24.5 23.1 19.7 18.3 21.2 17.5 16.8 22.4 20.6 23.9
 22. 11.9]
```

In [6]:

```
boston_features = data.feature_names
print(boston_features)
```

```
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DTS' 'RAD' 'TAX' 'PTRATIO']
```

```
['B' 'LSTAT']
```

In [7]:

```
# Preparing the Dataset
```

```
boston_df = pd.DataFrame(data.data, columns=boston_features)
boston_df.head()
```

Out[7]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

Here we can see that the target variable is not added, so we'll add the target column in the below steps

In [8]:

```
boston_df['PRICE']=data.target
```

Checking the newly added column in the dataframe

In [9]:

```
boston_df.head()
```

Out[9]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	PRICE
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2

In [10]:

```
boston_df.describe(include='all')
```

Out[10]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	PRICE
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	9.549407	408.2371	18.454382	392.8379	4.703158	32.063625
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707259	168.5371	2.839001	79.568379	0.861975	5.816131
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.0000	15.300000	392.8300	4.030000	21.600000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	279.0000	17.800000	396.9000	4.980000	24.000000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	5.000000	330.0000	18.700000	394.6300	5.330000	33.400000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	666.0000	20.300000	400.0000	6.575000	36.200000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.0000	20.300000	400.0000	9.140000	36.200000

In [11]:

```
boston_df.describe()
```

Out[11]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	9.549407	408.2371
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707259	168.5371
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.0000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	279.0000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	5.000000	330.0000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	666.0000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.0000

Checking the Transpose of the overall description of the dataset

In [12]:

```
boston_df.describe().T
```

Out[12]:

	count	mean	std	min	25%	50%	75%	max
CRIM	506.0	3.613524	8.601545	0.00632	0.082045	0.25651	3.677083	88.9762
ZN	506.0	11.363636	23.322453	0.00000	0.000000	0.00000	12.500000	100.0000
INDUS	506.0	11.136779	6.860353	0.46000	5.190000	9.69000	18.100000	27.7400
CHAS	506.0	0.069170	0.253994	0.00000	0.000000	0.00000	0.000000	1.0000
NOX	506.0	0.554695	0.115878	0.38500	0.449000	0.53800	0.624000	0.8710
...
TAX	506.0	408.237154	168.537116	187.00000	279.000000	330.00000	666.000000	711.0000
PTRATIO	506.0	18.455534	2.164946	12.60000	17.400000	19.05000	20.200000	22.0000
B	506.0	356.674032	91.294864	0.32000	375.377500	391.44000	396.225000	396.9000
LSTAT	506.0	12.653063	7.141062	1.73000	6.950000	11.36000	16.955000	37.9700
PRICE	506.0	22.532806	9.197104	5.00000	17.025000	21.20000	25.000000	50.0000

Checking the data type of each column variable

In [13]:

```
boston_df.dtypes
```

Out[13]:

CRIM	float64
ZN	float64
INDUS	float64
CHAS	float64
NOX	float64
RM	float64
AGE	float64
DIS	float64
RAD	float64
TAX	float64
PTRATIO	float64

```
PRICE      float64
B          float64
LSTAT      float64
dtype: object
```

In [14]:

```
boston_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  --
 0   CRIM        506 non-null   float64
 1   ZN          506 non-null   float64
 2   INDUS       506 non-null   float64
 3   CHAS        506 non-null   float64
 4   NOX         506 non-null   float64
 5   RM          506 non-null   float64
 6   AGE         506 non-null   float64
 7   DIS         506 non-null   float64
 8   RAD         506 non-null   float64
 9   TAX         506 non-null   float64
10  PTRATIO     506 non-null   float64
11  B           506 non-null   float64
12  LSTAT       506 non-null   float64
13  PRICE       506 non-null   float64
dtypes: float64(14)
memory usage: 55.5 KB
```

Checking whether any null value is present in the dataset

In [15]:

```
boston_df.isnull().sum()
```

Out[15]:

```
CRIM      0
ZN         0
INDUS      0
CHAS       0
NOX        0
RM         0
AGE        0
DIS        0
RAD        0
TAX        0
PTRATIO    0
B          0
LSTAT      0
PRICE      0
dtype: int64
```

Finding the correlation of the datapoints in the dataset

In [16]:

```
boston_df.corr()
```

Out[16]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	-0.379670	0.625505	0.582764	0.289946
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	0.664408	-0.311948	-0.314563	-0.391679

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	
INDUS	0.406583	0.533828	1.000000	0.062938	0.763651	0.391676	0.644779	0.708027	0.595129	0.720760	0.383248	0.35
CHAS	-0.055892	0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	-0.099176	-0.007368	-0.035587	-0.121515	0.04
NOX	0.420972	0.516604	0.763651	0.091203	1.000000	0.302188	0.731470	0.769230	0.611441	0.668023	0.188933	0.38
...
TAX	0.582764	0.314563	0.720760	0.035587	0.668023	0.292048	0.506456	0.534432	0.910228	1.000000	0.460853	0.44
PTRATIO	0.289946	0.391679	0.383248	0.121515	0.188933	0.355501	0.261515	0.232471	0.464741	0.460853	1.000000	0.17
B	-0.385064	0.175520	0.356977	0.048788	0.380051	0.128069	0.273534	0.291512	0.444413	0.441808	0.177383	1.00
LSTAT	0.455621	0.412995	0.603800	0.053929	0.590879	0.613808	0.602339	0.496996	0.488676	0.543993	0.374044	0.36
PRICE	-0.388305	0.360445	0.483725	0.175260	0.427321	0.695360	0.376955	0.249929	0.381626	0.468536	0.507787	0.33

Checking the shape of the dataframe used


```
In [17]:
boston_df.shape


Out[17]:
(506, 14)
```

Plotting the Heatmap

```
In [18]:
sns.set(rc={'figure.figsize': (15,10)})
sns.heatmap(boston_df.corr(),annot=True)

Out[18]:
<AxesSubplot:>
```





Data Visualization - Performing Basic EDA operations


1. Pairplot

In [19]:

```
sns.pairplot(boston_df)
```

Out[19]:

<seaborn.axisgrid.PairGrid at 0x17faabe61d0>




In [20]:


```
plt.scatter(boston_df['CRIM'],boston_df['PRICE'])
plt.xlabel("Crime Rate")
plt.ylabel("Price")
```

Out[20]:

Text(0, 0.5, 'Price')




2. Scatterplot - CRIME RATE vs PRICE


In [21]:

```
sns.regplot(x="CRIM",y="PRICE",data=boston_df)
```

Out[21]:

<AxesSubplot:xlabel='CRIM', ylabel='PRICE'>






3. ScatterPlot - AGE vs PRICE

In [22]:

```
plt.scatter(boston_df['AGE'], boston_df['PRICE'])  
plt.xlabel("Age")  
plt.ylabel("Price")
```

Out[22]:

Text(0, 0.5, 'Price')




In [23]:

```
sns.regplot(x="AGE", y="PRICE", data=boston_df)
```

Out[23]:

<AxesSubplot: xlabel='AGE', ylabel='PRICE'>






4. ScatterPlot - ROOMS vs PRICE

In [24]:

```
plt.scatter(boston_df['RM'],boston_df['PRICE'])  
plt.xlabel("Rooms Per Dwelling")  
plt.ylabel("Price")
```

Out[24]:

Text(0, 0.5, 'Price')




In [25]:

```
sns.regplot(x="RM",y="PRICE",data=boston_df)
```

Out[25]:

```
<AxesSubplot:xlabel='RM', ylabel='PRICE'>
```




In [26]:

```
sns.regplot(x="LSTAT",y="PRICE",data=boston_df)
```

Out[26]:

```
<AxesSubplot:xlabel='LSTAT', ylabel='PRICE'>
```



Detecting the outliers present in the Criminal Rate Data Column using Box-Plot

In [27]:


```
sns.set(rc={'figure.figsize':(10,10)})
sns.boxplot(boston_df['CRIM'])
```

C:\Users\Shobhandeb\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[27]:

<AxesSubplot:xlabel='CRIM'>



Detecting the outliers present in the AGE Data Column using Box-Plot

In [28]:


```
sns.set(rc={'figure.figsize':(10,10)})
sns.boxplot(boston_df['AGE'])
```

C:\Users\Shobhandeb\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[28]:

<AxesSubplot:xlabel='AGE'>



Detecting the outliers present in the PRICE Column using Box-Plot

In [29]:


```
sns.set(rc={'figure.figsize':(10,10)})
sns.boxplot(boston_df['PRICE'])
```


C:\Users\Shobhandeb\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[29]:

<AxesSubplot:xlabel='PRICE'>





Visualizing the DataFrame at a Glance using Pandas .head function

In [30]:

```
boston_df.head()
```

Out[30]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	PRICE
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2

Splitting the Dataset into Features and Label

In [31]:

```
X = boston_df.iloc[:, :-1]
Y = boston_df.iloc[:, -1]
```

In [32]:

```
X
```

Out[32]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33
...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.67
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.08
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.64
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.48
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.88

In [33]:

```
Y
```

Out[33]:

```
0      24.0
1      21.6
2      34.7
3      33.4
4      36.2
...
501    22.4
502    20.6
503    23.9
504    22.0
505    11.9
Name: PRICE, Length: 506, dtype: float64
```

Model Training

In [34]:

```
from sklearn.model_selection import train_test_split
```

Splitting the dataset into Train and Test Data into 70% and 30% Respectively

In [35]:

```
X_train,X_test,y_train,y_test = train_test_split(X,Y,test_size=0.3, random_state=15)
```

In [36]:

```
X_train
```

Out[36]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
129	0.88125	0.0	21.89	0.0	0.624	5.637	94.7	1.9799	4.0	437.0	21.2	396.90	18.34
349	0.02899	40.0	1.25	0.0	0.429	6.939	34.5	8.7921	1.0	335.0	19.7	389.85	5.89
257	0.61154	20.0	3.97	0.0	0.647	8.704	86.9	1.8010	5.0	264.0	13.0	389.70	5.12
60	0.14932	25.0	5.13	0.0	0.453	5.741	66.2	7.2254	8.0	284.0	19.7	395.11	13.15
314	0.36920	0.0	9.90	0.0	0.544	6.567	87.3	3.6023	4.0	304.0	18.4	395.69	9.28
...
375	19.60910	0.0	18.10	0.0	0.671	7.313	97.9	1.3163	24.0	666.0	20.2	396.90	13.44
133	0.32982	0.0	21.89	0.0	0.624	5.822	95.4	2.4699	4.0	437.0	21.2	388.69	15.03
396	5.87205	0.0	18.10	0.0	0.693	6.405	96.0	1.6768	24.0	666.0	20.2	396.90	19.37
245	0.19133	22.0	5.86	0.0	0.431	5.605	70.2	7.9549	7.0	330.0	19.1	389.13	18.46
456	4.66883	0.0	18.10	0.0	0.713	5.976	87.9	2.5806	24.0	666.0	20.2	10.48	19.01

In [37]:

```
X_train.shape
```

Out[37]:

```
(354, 13)
```

In [38]:

```
y_train.shape
```

```
Out[38]:
```

```
(354,)
```

```
In [39]:
```

```
y_test.shape
```

```
Out[39]:
```

```
(152,)
```

```
In [40]:
```

```
X_test.shape
```

```
Out[40]:
```

```
(152, 13)
```

Scaling the Train and Test Data using sklearn StandardScaler

```
In [41]:
```

```
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()
```

```
In [42]:
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
In [43]:
```

```
X_test_scaled = scaler.fit_transform(X_test)
```

```
In [44]:
```

```
X_train_scaled
```

```
Out[44]:
```

```
array([[ -0.32742136, -0.47138865,  1.5567594 , ...,  1.25320943,  
         0.45613396,  0.81342951],  
       [-0.43580738,  1.25986886, -1.45378412, ...,  0.56369103,  
         0.38227796, -0.93425629],  
       [-0.36172168,  0.3942401 , -1.05704583, ..., -2.51615782,  
         0.38070656, -1.0423461 ],  
       ...,  
       [ 0.30728278, -0.47138865,  1.00395126, ...,  0.7935305 ,  
         0.45613396,  0.95801718],  
       [-0.41516181,  0.48080298, -0.78137106, ...,  0.28788367,  
         0.37473522,  0.83027468],  
       [ 0.15426348, -0.47138865,  1.00395126, ...,  0.7935305 ,  
       -3.59201284,  0.90748168]])
```

```
In [45]:
```

```
X_test_scaled
```

```
Out[45]:
```

```
array([[ -0.39138546,  0.90829998, -0.71005094, ..., -1.08471245,  
         0.39241368, -0.47595831],  
       [-0.34329631,  0.31784256, -1.01978929, ..., -2.5383999 ,  
         0.2811574 , -0.97794122],  
       [-0.38109148, -0.52566804, -1.00810105, ..., -0.85024673,  
         0.40680566,  0.24974956],  
       ...,  
       [ 1.17984502, -0.52566804,  1.04464602, ...,  0.83790643,  
         0.40680566,  1.14325115],  
       [-0.37791113, -0.52566804, -0.34917654, ...,  1.16615844,  
         0.39141125,  0.80406711]])
```

```
0.3914125, 0.80486714],  
[-0.34087782, -0.52566804, 1.59837637, ..., 1.30683787,  
0.40680566, 0.77690152]])
```

In [46]:

```
from sklearn.linear_model import LinearRegression  
LR = LinearRegression()
```

In [47]:

```
LR
```

Out[47]:

```
▼ LinearRegression  
LinearRegression()
```

In [48]:

```
LR.fit(X_train,y_train)
```

Out[48]:

```
▼ LinearRegression  
LinearRegression()
```

In [49]:

```
## print the coefficients and the intercept  
print(LR.coef_)
```

```
[-7.30973225e-02  6.66062943e-02  8.45497046e-02  2.21512330e+00  
 -2.27372067e+01  3.24861978e+00  2.06578129e-02 -1.59247039e+00  
  3.48847293e-01 -1.39796398e-02 -9.58296625e-01  9.98858984e-03  
 -5.92254599e-01]
```

In [50]:

```
print(LR.intercept_)
```

```
41.76845495082349
```

In [51]:

```
## Prediction for the test data  
reg_pred=LR.predict(X_test)
```

In [52]:

```
reg_pred
```

Out[52]:

```
array([28.93841071, 40.17469652, 23.26283893, 22.72011976, 26.33677317,  
       6.50809139, 16.72675328, 13.83049735, 28.38006838, 16.83901688,  
       17.50579197, 22.45848043, 15.59048086, 16.11229233, 20.62101705,  
       15.20710548,  8.47374859,  7.69857378, 21.45782622, 10.97606569,  
       38.72583349, 13.26023439, 23.33227986, 19.27402726, 19.3360351 ,  
       19.62525449, 27.32359007, 19.91480848, 19.97039516, 19.98919575,  
       21.45883975,  7.54689782, 20.33795817, 19.38369205, 23.37468039,  
       19.05153146, 24.46267997, 28.19200979, 20.69966547, 18.68680301,  
       28.11584489, 35.29854655, 20.0879725 , 27.8604335 , 25.57788978,  
       21.59692292, 21.74601139, 30.24313863, 25.66136714, 20.36289475,  
       31.39205843, 15.24938636, 14.28689956, 14.33724217, 17.70236617,  
       30.67294605,  8.45801637, 29.38244272, 16.52514507, 26.35269311,  
       17.64563127, 27.64146931, 18.83849367, 30.29337701, 34.33685682,  
       20.40053045, 23.50840914, 18.29950906, 25.26796658, 18.96662268,  
       19.54785317, 23.47174606, 20.06513963, 12.71104176, 34.58884204,  
       20.18038844, 36.73034446, 17.86919358, 20.77346653, 13.21836627,
```


```
15.31962116, 17.54865527, 36.91230929, 22.57139152, 34.90624299,
18.75605528, 29.72788323, 23.65381032, 26.57590857, 20.73578394,
30.18801597, 24.57127722, 36.24961976, 23.92962871, 26.06991012,
32.79245337, 15.70279192, 27.04841957, 22.23639227, 22.14422298,
20.04264192, 12.51038729, 30.05440823, 3.91508266, 24.43278658,
13.96740998, 23.30762929, 24.83459343, 21.47895717, 35.16407276,
29.91813952, 23.00538811, 20.21991859, 26.7787996 , 18.54381139,
25.54068621, 29.75000896, 15.542685 , 15.20621604, 31.59691836,
24.48797917, 12.79369385, 20.42987578, 33.54619872, 30.10169183,
16.38623272, 22.77814 , 41.41331714, 34.03046388, 24.20339851,
34.37334928, 22.64842338, 8.61530203, 15.61618271, 13.20934389,
22.16702318, 32.47200523, 9.2780732 , 31.05708968, 27.49363576,
22.81749188, 17.4918494 , 20.47648792, 28.54144204, 23.17545754,
20.19885139, 19.7480775 , 17.29339103, 33.25948721, 18.95593496,
17.24115455, 16.74058242])
```

In [53]:

```
plt.scatter(y_test, reg_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")
```

Out[53]:

Text(0, 0.5, 'Test Predicted Data')



In [54]:


```
## residuals
residuals=y_test-reg_pred
```

In [55]:

```
sns.displot(residuals,kind="kde")
```

Out[55]:

<seaborn.axisgrid.FacetGrid at 0x17fb65292a0>




In [56]:

```
plt.scatter(reg_pred, residuals)
```

Out[56]:

<matplotlib.collections.PathCollection at 0x17fb8c36b30>



In [57]:

```
## Performance Metrics  
from sklearn.metrics import mean_squared_error    ## MSE
```

```
from sklearn.metrics import mean_absolute_error  ## MAE
print(mean_squared_error(y_test, reg_pred))
print(mean_absolute_error(y_test, reg_pred))
print(np.sqrt(mean_squared_error(y_test, reg_pred)))
```

```
27.58101431810535
3.7410552297491826
5.251762972384164
```

In [58]:

```
from sklearn.metrics import r2_score
score=r2_score(y_test, reg_pred)
print(score)
```

```
0.6745262607211457
```

In [59]:

```
## Adjusted R square
#display adjusted R-squared
1 - (1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[59]:

```
0.6438656910789349
```

Replacing the Xtrain and Xtest Data with their scaled value and check whether there is any significant changes in the R²

In [60]:

```
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
```

In [61]:

```
regressor.fit(X_train_scaled, y_train)
```

Out[61]:

```
▼ LinearRegression
LinearRegression()
```

In [62]:

```
print(regressor.coef_)
```

```
[-0.57477822  1.53891131  0.57966473  0.52327069 -2.61571068  2.30712032
  0.58203021 -3.24197468  3.03389615 -2.37665943 -2.0847086  0.95347108
 -4.21904768]
```

In [63]:

```
print(regressor.intercept_)
```

```
22.598870056497184
```

In [64]:

```
## Prediction for the test data
reg_pred_scaled=regressor.predict(X_test_scaled)
reg_pred_scaled
```

Out[64]:

```
array([29.42463134, 40.61980423, 23.29897949, 22.74061402, 26.4525965 ,
        6.28990307, 16.47974461, 13.56370461, 28.95497615, 17.49308995,
       17.47719586, 22.97436652, 16.3614426 , 15.45013899, 21.12309223,
       14.91234242,  8.43924416,  8.04022778, 21.83340869, 11.56712765,
       20.00751175, 12.22528082, 22.51226102, 10.21055180, 10.65126062])
```

```

33.03734473, 13.33320002, 23.31220433, 13.21033103, 13.03130303,
19.92020457, 28.03197224, 19.32099938, 19.96826541, 20.86872688,
22.19425428, 7.53319237, 20.773365, 19.39044594, 23.72136932,
18.87022269, 24.96021168, 28.23281392, 20.69551791, 19.37194935,
29.28811809, 35.34873981, 19.99782097, 28.54655757, 25.61194664,
22.23749836, 21.74517391, 30.55628098, 26.44523454, 20.37401324,
32.33920514, 14.87300921, 14.39367596, 14.45346953, 17.92770041,
31.34936014, 8.73807517, 29.57151886, 16.30931465, 26.45033129,
17.47457576, 27.75310272, 18.71862476, 30.21772465, 34.66598234,
20.40381065, 23.27971114, 18.24107235, 25.64969651, 18.94833594,
20.15482152, 24.18153851, 20.07589036, 12.03550143, 35.02016365,
21.25567391, 37.04163541, 18.4118881, 20.74489414, 13.39241986,
15.23374827, 17.61224194, 37.55997143, 23.70708553, 35.31240389,
18.59071991, 30.55847821, 23.6698027, 26.6435274, 20.9614213,
30.74266834, 25.44464509, 35.99961779, 24.06620718, 25.83510891,
32.60717018, 16.1425784, 27.7594317, 22.0735267, 21.78537948,
20.03481481, 12.71558176, 30.65271935, 3.36324905, 24.63843861,
13.96277842, 23.67535194, 24.59990408, 21.54826575, 35.59706163,
30.12203431, 23.03074256, 20.97569756, 26.07544888, 19.00141849,
25.34055397, 29.91583008, 15.45227059, 14.32601875, 32.36723659,
24.67705534, 12.18934835, 20.32484065, 32.96968048, 30.56629772,
15.94230652, 22.82085236, 41.36355971, 34.37172257, 24.61744623,
34.76112273, 22.98704168, 8.57551228, 15.90253854, 12.85275527,
22.37127071, 33.12591892, 9.33521016, 31.56154069, 28.16215999,
23.35092939, 17.32183768, 20.98538519, 28.98690929, 23.22604912,
20.98033326, 19.73015481, 18.5988905, 33.46069977, 19.05492711,
17.1429949, 16.62854312])

```

In [65]:


```

plt.scatter(y_test, reg_pred_scaled)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")

```

Out[65]:

Text(0, 0.5, 'Test Predicted Data')



In [66]:

```
## residuals
residuals_scaled=y_test-reg_pred_scaled
```

In [67]:

```
residuals_scaled
```

Out[67]:


```
301    -7.424631
262     8.180196
172    -0.198979
505   -10.840614
111    -3.652596
...
380    -8.198891
307    -5.260700
381    -8.154927
106     2.357005
139     1.171457
Name: PRICE, Length: 152, dtype: float64
```

In [68]:

```
sns.displot(residuals_scaled,kind="kde")
```

Out[68]:

<seaborn.axisgrid.FacetGrid at 0x17fb8c82f50>





In [69]:

```
plt.scatter(reg_pred,residuals_scaled)
```

Out[69]:

<matplotlib.collections.PathCollection at 0x17fb8d721a0>





In [70]:

```
## Performance Metrics
from sklearn.metrics import mean_squared_error    ## MSE
from sklearn.metrics import mean_absolute_error    ## MAE
print(mean_squared_error(y_test, reg_pred_scaled))
print(mean_absolute_error(y_test, reg_pred_scaled))
print(np.sqrt(mean_squared_error(y_test, reg_pred_scaled)))
```

```
27.46473809683593
3.727436748688041
5.2406810718489565
```

In [71]:

```
from sklearn.metrics import r2_score
score_scaled=r2_score(y_test, reg_pred_scaled)
print(score_scaled)
```

```
0.6758983950483787
```

In [72]:

```
## Adjusted R square
#display adjusted R-squared
1 - (1-score_scaled)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[72]:

```
0.6453670844369941
```

There is no such significant changes between the Normal and Scaled R^2 Values

Evaluating the Dataset using Ridge Regression

In [73]:

```
## Ridge
from sklearn.linear_model import Ridge
ridge=Ridge()
```

In [74]:

```
ridge.fit(X_train, y_train)
```

Out[74]:

▼ Ridge
Ridge()

In [75]:

```
ridge_pred = ridge.predict(X_test)
```

In [76]:

```
ridge_pred
```

Out[76]:


```
array([[28.34700411, 40.99020024, 22.61027099, 23.2155383 , 26.11759627,
        5.98219852, 15.62597673, 12.83874366, 27.99127334, 16.39679102,
        17.63269965, 22.29196843, 15.78474533, 16.47304297, 20.43928823,
        15.53700685,  8.49051102,  7.8570057 , 22.21636121, 10.37300533,
        39.15601791, 13.50243764, 22.42297964, 19.37771 , 20.02993273,
        19.47530705, 26.96124868, 20.72561217, 20.17750241, 20.71514059,
        21.94343695,  9.32729636, 20.38013688, 19.1333901 , 23.19582726,
        20.02576519, 23.47458314, 28.86578648, 20.87954425, 18.47360772,
        28.94436027, 34.87608624, 20.54470438, 27.33611564, 25.40585055,
        21.40416119, 21.0462387 , 30.30832271, 26.11474415, 20.26341559,
        31.91067008, 17.09132649, 14.36271739, 14.53180137, 17.96166841,
        31.01618851,  8.82019251, 28.42608421, 15.4304323 , 26.85162603,
        16.61153694, 27.46829035, 18.41043268, 29.5195852 , 34.09605507,
        20.22223502, 22.90603938, 18.54047753, 25.00272711, 19.63226713,
        18.43735905, 23.78779977, 20.24490781, 13.18344169, 34.96474836,
        20.79023033, 37.42422335, 17.13861604, 21.01270366, 13.99123209,
        14.4508423 , 17.43909906, 36.97607769, 22.90718041, 35.16515341,
        19.2989718 , 29.88230126, 23.07193895, 26.22457409, 20.75840548,
        30.13141762, 24.5859121 , 37.01350018, 23.78615258, 25.57973946,
        32.02685983, 15.78405085, 27.41134644, 21.26214375, 23.12465509,
        19.41357858, 12.66093485, 29.90958765,  4.38255416, 24.25073335,
        15.91089374, 22.97921236, 23.89152687, 20.77540019, 34.81754799,
        30.12105276, 22.52794502, 19.86807171, 26.46315873, 18.84789258,
        24.91207356, 29.2459565 , 16.2150825 , 14.22966397, 31.55193231,
        24.08301293, 12.64748472, 19.93409229, 32.92761943, 29.70919044,
        16.78158311, 22.41770084, 40.88693989, 34.39226363, 24.06135042,
        33.92040686, 22.83345534, 10.45132671, 15.85212532, 14.81175304,
        22.48341944, 32.44637478,  9.10872096, 30.80920283, 27.10737493,
        22.70310752, 16.47854624, 20.17080451, 28.16656697, 22.97102641,
        20.38098202, 20.3644423 , 18.08376687, 33.00326745, 18.90459902,
        16.85547231, 16.76993899])
```


In [77]:

```
plt.scatter(y_test,ridge_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")
```

Out[77]:

```
Text(0, 0.5, 'Test Predicted Data')
```





In [78]:

```
## residuals
ridge_residuals=y_test-ridge_pred
ridge_residuals
```

Out[78]:


```
301    -6.347004
262     7.809800
172     0.489729
505   -11.315538
111    -3.317596
...
380    -7.683767
307    -4.803267
381    -8.004599
106     2.644528
139     1.030061
Name: PRICE, Length: 152, dtype: float64
```

In [79]:

```
sns.displot(ridge_residuals,kind="kde")
```

Out[79]:

<seaborn.axisgrid.FacetGrid at 0x17fb8d9f850>




In [80]:

```
plt.scatter(ridge_pred,ridge_residuals)
```

Out[80]:

<matplotlib.collections.PathCollection at 0x17fb8ea9540>



In [81]:

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,ridge_pred))
print(mean_absolute_error(y_test,ridge_pred))
print(np.sqrt(mean_squared_error(y_test,ridge_pred)))
```

26.895629415693488
3.6794629931440492
5.186099634184971

In [82]:

```
from sklearn.metrics import r2_score
ridge_score=r2_score(y_test,ridge_pred)
print(ridge_score)
```

0.6826142441600589

In [83]:

```
## Adjusted R square
#display adjusted R-squared
1 - (1-ridge_score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[83]:

0.6527155860012239

Evaluating the Dataset using the Lasso Regression

In [84]:

```
from sklearn.linear_model import Lasso
lasso = Lasso()
```

In [85]:

lasso

Out[85]:

- ▼ Lasso

In [86]:

```
lasso.fit(X_train, y_train)
```

Out[86]:

- ▼ Lasso

In [87]:

```
print(lasso.intercept )
```

42.22410959384636

In [88]:

```
print(lasso.coef )
```

```
[-0.02663527  0.07056985 -0.          0.          -0.          0.5581624
  0.03947285 -0.70744747  0.26299305 -0.01560035 -0.68467839  0.00800488
 -0.81411349]
```

In [89]:

```
lasso_pred=lasso.predict(X_test)
```

In [90]:

lasso pred

Out[90]:

```
array([27.12941361, 36.88288348, 23.53663729, 25.39359024, 24.84286728,
        4.74509463, 13.77197408, 14.59393706, 27.91697514, 20.70575638,
        17.19959477, 23.38059776, 18.76539026, 15.71949933, 21.9400187 ,
        16.89252795,  7.77093828,  7.50463158, 23.9060242 , 10.67279224,
        36.68242587, 16.43054605, 21.60211303, 18.79002068, 22.97393494,
        21.98666538, 27.36347518, 20.51518792, 20.88869769, 24.0578185 ,
        24.76600528, 11.71271316, 21.09083519, 18.65597701, 22.96610777,
        24.27196609, 23.65172725, 30.72143679, 19.57450749, 20.41043529,
        26.31788589, 30.87741489, 20.55445125, 28.70729577, 22.47098468,
        23.75433348, 21.10560936, 28.90566851, 28.19825584, 16.67068047,
        32.12026299, 20.94566571, 16.26353947, 16.30748208, 19.98072517,
        30.37613072,  8.87735163, 26.18796243, 12.67216926, 27.34816607,
        17.23128131, 27.34355876, 18.57341514, 30.83959245, 33.33076023,
        20.40398752, 23.76843679, 19.73860467, 26.12468643, 19.82619535,
        19.14066668, 24.65846193, 19.44224362, 14.85090005, 34.12218508,
        24.57761427, 35.32778313, 17.67055649, 20.75709165, 14.43617187,
        13.07633825, 15.7003081 , 34.14255471, 24.5261761 , 31.5570532 ,
        18.31798381, 29.61277049, 23.18107675, 26.92377089, 23.20020534])
```

```

29.55871397, 27.93439204, 29.58002763, 24.96195126, 23.44437412,
28.19490581, 17.65750882, 27.42232493, 19.52255178, 21.93159263,
18.13777849, 11.19682599, 29.79131801, 2.30146038, 24.03640166,
18.25897984, 22.23769927, 21.57007722, 20.81030838, 30.89776734,
29.08864257, 21.89925985, 23.03711848, 25.66765445, 19.68908392,
23.18863185, 28.70904018, 18.68770727, 14.57359524, 30.66186561,
24.58434053, 10.89134835, 19.84919038, 30.72379031, 30.70017339,
18.0433517 , 22.58478993, 34.40515005, 30.02136614, 23.9168915 ,
31.20069332, 24.54281215, 12.44231477, 17.88352247, 14.46959123,
22.73879405, 31.34618692, 5.81778454, 31.55147121, 28.4583525 ,
24.25812821, 14.77655907, 21.84372431, 26.78499449, 22.59591976,
22.40993065, 19.96668011, 17.62667583, 31.72425704, 16.39858989,
16.53095951, 16.19477664])

```

In [91]:


```

plt.scatter(y_test,lasso_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")

```

Out[91]:

Text(0, 0.5, 'Test Predicted Data')



In [92]:

```

## residuals
lasso_residuals=y_test-lasso_pred
lasso_residuals

```

Out[92]:

```

301    -5.129414
262     11.917117
172     -0.436637
505    -13.493590

```




```
111      -2.042867
      ....
380     -7.226676
307     -3.524257
381     -5.498590
106      2.969040
139      1.605223
Name: PRICE, Length: 152, dtype: float64
```

In [93]:

```
sns.displot(lasso_residuals,kind="kde")
```

Out[93]:

<seaborn.axisgrid.FacetGrid at 0x17fb8eed330>





In [94]:

```
## SCatter plot with predictions and residual
##uniform distribution
plt.scatter(lasso_pred,lasso_residuals)
```

Out[94]:

<matplotlib.collections.PathCollection at 0x17fba74c760>





In [95]:

```
sns.regplot(lasso_pred,lasso_residuals)
```

C:\Users\Shobhandeb\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[95]:

<AxesSubplot:ylabel='PRICE'>



In [96]:

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,lasso_pred))
print(mean_absolute_error(y_test,lasso_pred))
print(np.sqrt(mean_squared_error(y_test,lasso_pred)))
```

33.31417646922572
4.044140676389867
5.771843420366298

In [97]:

```
from sklearn.metrics import r2_score
lasso_score=r2_score(y_test,lasso_pred)
print(lasso_score)
```

0.6068712534869765

In [98]:

```
## Adjusted R square
#display adjusted R-squared
1 - (1-lasso_score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[98]:

0.5698373860618365

Evaluating the ElasticNet Regression

In [99]:

```
from sklearn.linear_model import ElasticNet
elasticnet = ElasticNet()
```

In [100]:

elasticnet

Out[100]:

▼ ElasticNet
ElasticNet()

In [101]:

```
elasticnet.fit(X_train,y_train)
```

Out[101]:

▼ ElasticNet
ElasticNet()

In [102]:

```
print(elasticnet.intercept_)
```

41.43181147909337

In [103]:

```
print(elasticnet.coef_)
```

```
[-0.04587096  0.07269826 -0.          0.          -0.          0.79596544
  0.03686871 -0.7795076   0.29552182 -0.01671115 -0.70696726  0.00824164
 -0.79346039]
```

In [104]:

```
elasticnet_pred=elasticnet.predict(X_test)
```

In [105]:

elasticnet_pred

Out[105]:

```
array([27.26828373, 37.51032106, 23.5103766 , 25.16133916, 24.8620994 ,
        4.96350039, 13.75939807, 14.15667878, 27.93675223, 20.29567914,
```

```

1.36336833, 13.7333307, 11.1300707, 27.3307323, 20.2330731,
17.11111674, 23.24198129, 18.29521516, 15.94525879, 21.63377027,
16.40644521, 7.87286401, 7.22000653, 23.84317529, 10.73149498,
37.09206294, 16.14321934, 21.46168756, 18.71686962, 22.66965331,
21.75464982, 27.23463365, 20.62922864, 20.67907094, 23.58241666,
24.58592064, 11.78176239, 20.80240837, 18.79055074, 22.951107 ,
24.17318451, 23.54136541, 30.86092198, 19.808766 , 20.29363844,
26.46814145, 31.09658387, 20.6554074 , 28.63220283, 22.81079176,
23.5971913 , 21.19882451, 28.96528236, 28.01436765, 17.07321262,
32.26511008, 20.85186443, 16.06946393, 16.03035783, 19.76871859,
30.34973931, 8.49742367, 26.34728922, 12.96971314, 27.25290027,
17.19984168, 27.43741824, 18.47001984, 30.68086412, 33.4438989 ,
20.28502535, 23.64998184, 19.64133836, 25.98274388, 20.02589897,
18.98966344, 24.53715483, 19.41113857, 14.84442847, 34.27319022,
24.07758544, 35.79058396, 17.59623148, 20.89894275, 14.3681251 ,
13.31437778, 15.81759266, 34.35110555, 24.3429957 , 31.92712315,
18.59530486, 29.55954147, 23.26274262, 26.90558219, 22.98676213,
29.53320464, 27.56199386, 30.2395347 , 24.87823579, 23.54384238,
28.43667746, 17.27931099, 27.3885288 , 19.4826149 , 22.12587463,
18.26416243, 11.20908887, 29.71976701, 2.60692363, 24.14034378,
18.30165931, 22.13797235, 21.69294449, 20.83264279, 31.11039637,
29.20630786, 21.89492996, 22.71580248, 25.65511909, 19.59974065,
23.21237664, 28.75998786, 18.43942393, 14.54893282, 30.72271777,
24.62367812, 11.14817335, 19.78183593, 30.81838114, 30.59386778,
17.97997007, 22.66230113, 34.72279392, 30.46533435, 23.87706966,
31.36784604, 24.37016055, 12.54069607, 17.67449394, 14.57784374,
22.73312952, 31.35300892, 6.12363363, 31.4891038 , 28.38592949,
24.17505825, 15.00700429, 21.5219948 , 26.93744597, 22.54457767,
22.23913207, 20.00321318, 16.47854978, 32.03802161, 16.60981382,
16.4738673 , 16.12693736])

```

In [106]:


```

plt.scatter(y_test,elasticnet_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")

```

Out[106]:

Text(0, 0.5, 'Test Predicted Data')



10 20 30 40 50
Test Truth Data

In [107]:


```
sns.regplot(y_test,elasticnet_pred)
```

C:\Users\Shobhandeb\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[107]:

<AxesSubplot:xlabel='PRICE'>



In [108]:

```
## residuals
elasticnet_residuals=y_test-elasticnet_pred
elasticnet_residuals
```

Out[108]:

```
301    -5.268284
262    11.289679
172    -0.410377
505   -13.261339
111    -2.062099
...
380    -6.078550
307    -3.838022
381    -5.709814
106     3.026133
139     1.673063
```


```
108 1090000
Name: PRICE, Length: 152, dtype: float64
```

```
In [109]:
```

```
sns.displot(elasticnet_residuals,kind="kde")
```

```
Out[109]:
```

```
<seaborn.axisgrid.FacetGrid at 0x17fba7ae2f0>
```




```
In [110]:
```

```
## Scatter plot with predictions and residual
##uniform distribution
plt.scatter(elasticnet_pred,elasticnet_residuals)
```

```
Out[110]:
```

```
<matplotlib.collections.PathCollection at 0x17fba83fb80>
```



-15

5

10

15

20

25

30

35

In [111]:


```
sns.regplot(elasticnet_pred,elasticnet_residuals)
```

C:\Users\Shobhandeb\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[111]:

<AxesSubplot:ylabel='PRICE'>



In [112]:

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,elasticnet_pred))
print(mean_absolute_error(y_test,elasticnet_pred))
print(np.sqrt(mean_squared_error(y_test,elasticnet_pred)))
```

32.36843282824359

3.9801854082590764

5.6893262191795255

In [113]:

```
from sklearn.metrics import r2_score
elasticnet_score=r2_score(y_test,elasticnet_pred)
print(elasticnet_score)
```

0.6180316377889994

In [114]:

```
## Adjusted R square  
#display adjusted R-squared  
1 - (1-elasticnet_score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[114]:

0.5820491109140501

In []: